

The Reverse Centaur's Guide to Life After AI

Cory Doctorow

Introduction

In May 2025, the internet briefly chose a new, unfortunate main character: Marco Buscaglia, a Chicago-based freelance writer behind a syndicated summer reading list. Though the list was credited to him, he wasn't its principal author. Rather, the list was written by a chatbot.

We know this is true for two reasons. First, because ten of the fifteen books on the list were nonexistent. They were AI “hallucinations” (a term that AI boosters use because it sounds sexier than “errors”). Second, because Buscaglia confessed to using a chatbot, in an interview with the veteran tech reporter Jason Koebler, who cofounded the journalist-owned news site 404 Media.

In Koebler's article, Buscaglia admits that he used a chatbot to source the list, and blames the inclusion of nonexistent books on the list on his own oversight: “I do use AI for background at times but always check out the material first. This time, I did not and I can't believe I missed it because it's so obvious.”

Buscaglia comes across as utterly mortified, telling Koebler, “No excuses. On me 100 percent and I'm completely embarrassed. It's a complete mistake on my part ... This is just idiotic of me, really embarrassed.”

For a few days, the internet had a lot of fun with poor Buscaglia, from the baffled authors whose nonexistent books were

recommended on his list to the legion of people who are sick to the back teeth of having AI jammed into their eyeballs by overcapitalized tech companies.

But Koebler's analysis went beyond Buscaglia's contrition. In follow-up stories and on the 404 Media podcast, Koebler dove deep into the context of Buscaglia's unfortunate moment in the internet's baleful glare.

Start with the list itself. The summer reading list was many lists and guides in a sixty-four-page "Best of Summer" supplement that was distributed to multiple daily newspapers by King Features Syndicate, a division of Hearst Communications, a massive and established publishing conglomerate.

Buscaglia wrote (or "wrote") the majority of those lists. This is a remarkable feat! Koebler has experience writing lists like these, having begun his career as an intern at *Washington Monthly* magazine, where a list like this would be assigned to a team of three interns, overseen by an experienced newsroom journalist and backstopped by a large and skilled fact-checking department.

I've done this kind of work myself. I got my start as a journalist with assignments from *Wired* magazine that had me devoting an entire month to evaluating dozens of multitools, speaking to experts and manufacturers' representatives, in order to pick my top three and write a brief paragraph about each. These were fun assignments, but they were also serious and involved endless back-and-forth with the fact-checkers, even down to providing proof of how many different blades and attachments each tool had.

To be frank, there's *no way* that Buscaglia could have written all those lists and summer guides (including a guide to hammocks that features a quote from a nonexistent professor of leisure studies who is credited with writing a scholarly journal article on "hammock culture"—that paper also doesn't exist).

In 2025, any editor who assigns a freelancer to "write" the bulk of a sixty-four-page summer guide is either tacitly or explicitly commissioning that writer to confine their efforts to prompting a chatbot to barf up a bunch of plausible verbiage and hammer it into shape.

Buscaglia was asked to do the work of dozens of writers, researchers, editors, and fact-checkers, on a short timescale that guaranteed that the resulting product would be riddled with AI-

generated errors. We don't know how much Buscaglia was paid for this job, but it's a sure bet that it doesn't add up to the total salary of all the skilled professionals whose jobs he was asked to do.

In other words, Buscaglia was set up to fail. His job wasn't merely to oversee a chatbot—it was to absorb the blame for that chatbot's mistakes.

Like Buscaglia, I am a freelance writer. What's more, I'm a freelance writer who has recently used AI in the course of my work. I was trying to locate a quote from an expert I'd heard say something very smart on a podcast. The only problem was that I couldn't remember the name of the expert *or* which podcast I'd heard him on.

So I downloaded an open-source AI speech transcription model called Whisper and installed it on my laptop. Then I downloaded thirty or forty hours' worth of podcasts that I'd recently listened to, and fed them all to Whisper, asking it to produce a transcript of all of them.

While Whisper chugged away in the background on my laptop, I kept working on the article, answering emails and alt-tabbing to social media in case someone had uploaded a video of a guy putting a lemon up his nose (normal writer stuff, in other words).

A couple hours later, I opened the folder where Whisper dumps its transcripts and searched them for some phrases I remembered from the quote I was looking for, and located the file in minutes (it was Hagen Blix, coauthor of *Why We Fear AI*, discussing the plight of AI therapists on the excellent *This Machine Kills* podcast). Then I used the time code Whisper provided to open the podcast file and clean up the transcribed quote, which I pasted into my article (I was right; it was a great quote).

This was a *great* experience. It was a great *AI* experience. Thanks to AI, my computer has now acquired a permanent new feature: it can turn arbitrary amounts of recorded speech into pretty reliable transcripts, in a manner so efficient that my laptop's fan doesn't even turn on. The fact that Whisper is open-source means that it can be maintained forever, by anyone in the world who has the skills and desire to do so, no matter what happens to Whisper's manufacturer (a grossly overhyped and terrible firm called OpenAI).

Two freelance writers, using AI. One was made miserable and embarrassed by his AI usage. The other had a delightful experience, saved a bunch of time, and produced a better piece of writing. How

to explain the paradox?

One possibility is that the difference lies in *how* we used AI. I used AI to transcribe some audio, whereas Buscaglia used it to generate some writing.

Though I don't brainstorm with AI or other forms of automated or random text generation, I don't have a moral objection to these practices. There are plenty of writers who find text generation to be a productive tool for producing excellent work. The Dadaists' "cut-ups" involved slicing phrases out of a manuscript and scrambling them around, looking for serendipitous juxtapositions. Writers delight in Brian Eno and Peter Schmidt's "Oblique Strategies" deck, which contains fifty-five cards printed with mysterious orders ("Be the first person to not do something that no one else has ever not done before"). The Surrealists loved their "Exquisite Corpse," in which a group of writers take it in turn to append text to the previous writer's text, but each writer can see only the last sentence of the work when they begin.

Today, I know lots of writers who use chatbots to produce fine work: they might use the chatbot as a sounding board for evaluating ideas (or variations on ideas), or to challenge them with writing prompts, or to suggest improvements. As I said, I don't use AI that way, but I know people who do, and I like the things they write.

The difference between Buscaglia's AI experience and my own—the difference between a useful tool and a technological torment—is the difference between a *centaur* and a *reverse centaur*.

In automation theory (the academic study of automation), a *centaur* is a person who is assisted *by* a machine. Think of a clever human's head, arms, and hands atop a horse's strong body. Riding a bicycle or driving a car makes you a centaur; so does using the mute button on your TV remote when an ad comes on. Wearing a hearing aid makes you a centaur, and so does using a calculator to multiply large numbers. Being a centaur can be glorious. My whole writing career has been a series of centaur moves, from the used IBM Selectric my parents bought for me to play with when I was six or seven to the Apple II Plus we got when I was nine, which kicked off an unbroken string of better and better writing tools, with spellcheckers, version control, collaboration features, autosave, and more. Best of all, I get to choose *exactly* which of these features I use. I work for myself, after all, so if I don't want to use the grammar checker (hell

no), that's my business. No one expects me to write more pages when I get a new tool. A couple of years ago, a former student of mine asked me to try out his LLM tool that would help me write dialogue and flesh out characters. I played with it for ten minutes and then never went back to it. No one told me I was being uncooperative or spoiling a grand plan to increase efficiency and realize cost savings by refusing to use AI. Sometimes, when I'm really stuck, I write with a pen in a notebook. No one cares, except me.

A *reverse centaur* is a human who is conscripted into acting as an assistant *to* a machine. There's a classic *I Love Lucy* episode where Lucy and Ethel are working on an assembly line at a chocolate factory, taking bonbons off the belt and wrapping them in paper. As the belt goes faster and faster, Lucy and Ethel have to work at superhuman speed. They're reverse centaurs: the machine can move the chocolates from one place to another, but it needs a human to pack them into the box, and the humans who act as its assistant are made to work at a pace that exceeds all human capacity until calamity ensues.

I Love Lucy played that bit for laughs. Amazon warehouse workers get the horror-movie version. They are observed by AI-equipped cameras that time their movements and monitor their "time off-task," penalizing them if they fail to make quota. Amazon warehouse workers experience the highest level of on-the-job injuries in the U.S. warehouse sector, and many of them have to resort to urinating in bottles because a visit to the toilet would blow their quotas. An Amazon warehouse is full of machines, but there are jobs the machines can't do, and that's where the Amazon warehouse workers come in—they assist the machines. They are reverse centaurs, and they have been conscripted to serve as peripherals for the warehouse's automation systems. They aren't merely used by those machines—they are used *up*.

It's not hard to imagine how a warehouse worker might *choose* to use AI in their daily work—for example, a computer vision system in a pair of smart glasses highlights an item they're looking for in a bin (there've been many instances where I've stared directly at something without seeing it in which *I* would have loved this). Automation isn't *necessarily* the enemy of warehouse work: there's nothing wrong with a forklift! The difference between automation that helps a warehouse worker and automation that torments that worker

is whether the worker gets to choose where, when, and how to use that automation. It's the difference between a centaur and a reverse centaur.

I love the automation system in my car that warns me if I'm drifting out of my lane, and I recently discovered (the hard way!) that if the person ahead of me brakes suddenly, my car will let out a sphincter-puckering series of beeps and activate its own brake. I was pretty happy about that, even if it did come as a hell of a surprise.

Compare that with the drivers in those Amazon vans rolling around your neighborhood. They have to sign into at least *nine* separate apps, and they are continuously scored based on their driving performance, as assessed by various AI tools. Drivers lose points for braking or swerving (even if that's the only way to avoid a sudden road hazard) or for deviating from the proscribed route set by the AI (even if there are obstructions or hazards). Drivers have impossible-to-meet quotas and the per-parcel compensation rate drops if they fail to meet it. Drivers are forbidden from peeing in bottles, but also given no time to urinate. The driver is just a peripheral for the van, present only because the van can't drive itself or get your parcel onto your porch. They are a reverse centaur.

This centaur/reverse centaur distinction is the heart of the paradox at the heart of the debate about the usefulness of AI tools. When you find yourself surrounded by people swearing that a given tool is worse than useless and others swearing that it has made their lives easier and better, you can bet that the former group is made up of reverse centaurs who've had AI imposed upon them, the latter group is all centaurs who've gotten to make up their own minds about where, when, and how to use AI tools. The solution to the paradox is to stop thinking about what the gadget does, and pay attention to who the gadget does it *to* and who the gadget does it *for*. The important part isn't the technical characteristics of the device, it's the power relationships of the people who use the device.

It's been a long time since I last held a factory job. I'm a science fiction writer by trade, which means that I get asked about AI about a million times per day. Some people think science fiction writers are seers or prophets (regrettably, some extremely delusional science fiction writers share this view). They imagine that science fiction is a literature that predicts the future.

This is nonsense, of course. The future isn't predictable, which is

a damned good thing, because if the future were predictable, it would be foreordained, which would mean that the actions taken by people don't matter.

But science fiction writers *do* have an intimate relationship with the future. By creating futuristic parables—made-up stories about futures that may never come to pass—we SF writers remind everyone that the future is up for grabs, full of possibility.

That's in stark contrast to tech bosses, who claim total authority over our possibilities, both present and future. There's Mark Zuckerberg: *Yes, well, obviously you'd like to talk to your friends without my spying on you all from asshole to appetite, but let's be realistic here. That's like asking for water that isn't wet.*

Or Apple CEO Tim Cook: *Yes, certainly, it would be nice if we could provide you with a reliable mobile device without locking you into our App Store, where we rake in a 30 percent commission on everything you spend, and where we get to decide which apps you can use, and which ones you can't (like the privacy apps that we banned in China or the ICE-tracking app we banned in America). But come on, it's just not realistic to want a computer that a) works and b) takes orders from you, not me.*

This is a cheap bully's trick: insisting that their abusive behavior is out of their hands, that they are merely acting in accord with some kind of iron law or great force of history. Margaret Thatcher was the world champion of this: she insisted that her cruel program of austerity and privation was inevitable, repeating the phrase "There is no alternative" so often that wags started referring to her as "Tina" ("There Is No Alternative").

Thatcher's contemporary successors are to be found in the C-suites of the tech industry.

Call this philosophy *inevitabilism*, the insistence that there is only one conceivable way to do things, and any problems you're experiencing aren't anyone's fault, they're just inescapable *reality*. The fact that the world is filled with sneaky Bluetooth sniffers that track your location as you walk through shopping centers or even just down the street isn't something you can blame a person or a company for; it's just a fact.

Science fiction is an anti-inevitabilist literature. Science fiction is only incidentally about thinking up new gadgets and explaining what they do. Far more important than what the gadget *does* is who it

does it *to* and who it does it *for*. As the legendary science fiction editor Gardner Dozois once wrote, “Most SF can predict the car, some SF can predict the drive-in theater, but SF that can predict the changes in teen-age sexual behavior as a result of the drive-in is vanishingly rare.” To this I would add, “Truly visionary science fiction might think about how a world where you need a government-issued ID like a driver’s license in order to engage in sexual experimentation might plausibly lead to a ‘database nation’ of ubiquitous surveillance.”

The mere existence of a literature in which many different writers have imagined many different futures, each of which can *feel* possible and even desirable, is a rebuke to inevitabilism. The fact that social relations between people and their technology can be different means that the current arrangement is a choice, and—crucially—it means that we can choose something else.

AI hucksters want you to believe that all the things they call “AI”—an incoherent grab bag of many technologies, some of them not especially related to the rest—are coming, and that when they arrive, there is only one conceivable way that we could use them. This is just high-tech Thatcherism, the inevitabilist move of a bully who insists that they’re only doing what implacable reality demands of them.

This is a book about what AI can and *cannot* do, but even more important, it’s about the possible social arrangements of AI, from not using some AI technology at all, to using it in ways that let some of us choose to be centaurs, while saving our friends and neighbors from being conscripted into reverse-centaury.

The current wave of AI is full of software performing impressive feats in both parsing and generating language and images and sounds. There are lots of interesting, fun, and productive ways to use this technology. There is nothing about the *technology* of AI that determines how it *must* be used. We can choose to use it sometimes, or never, or all the time, depending on our needs and proclivities. We don’t have to let billionaires tell us how it must be used.

Margaret Thatcher’s “There is no alternative” has a fine rejoinder in the science fiction writer William Gibson’s famous maxim, “The street finds its own use for things.”

SF writers make lousy prophets, but we can be pretty good technology critics.